

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



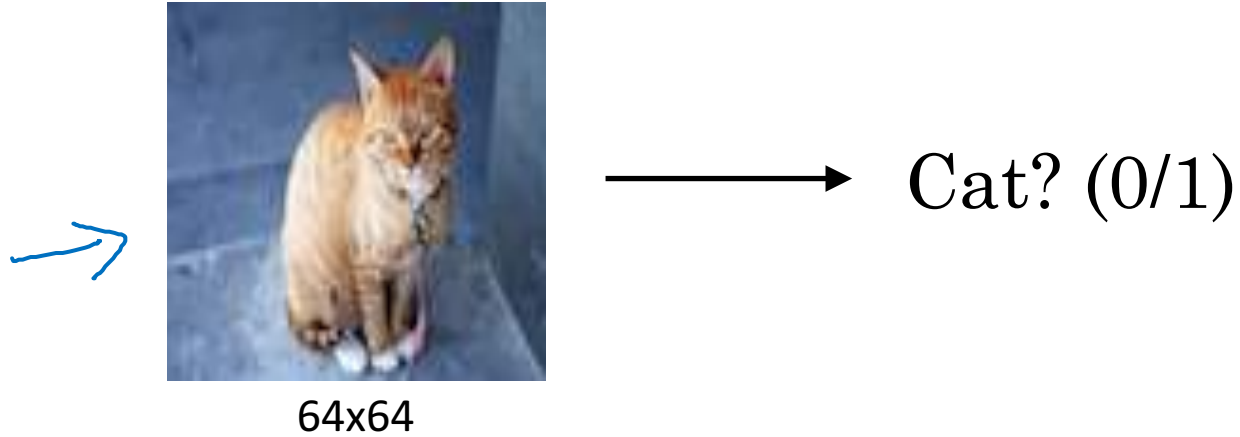
deeplearning.ai

Convolutional Neural Networks

Computer vision

Computer Vision Problems

Image Classification





Neural Style Transfer



Object detection



Deep Learning on large images



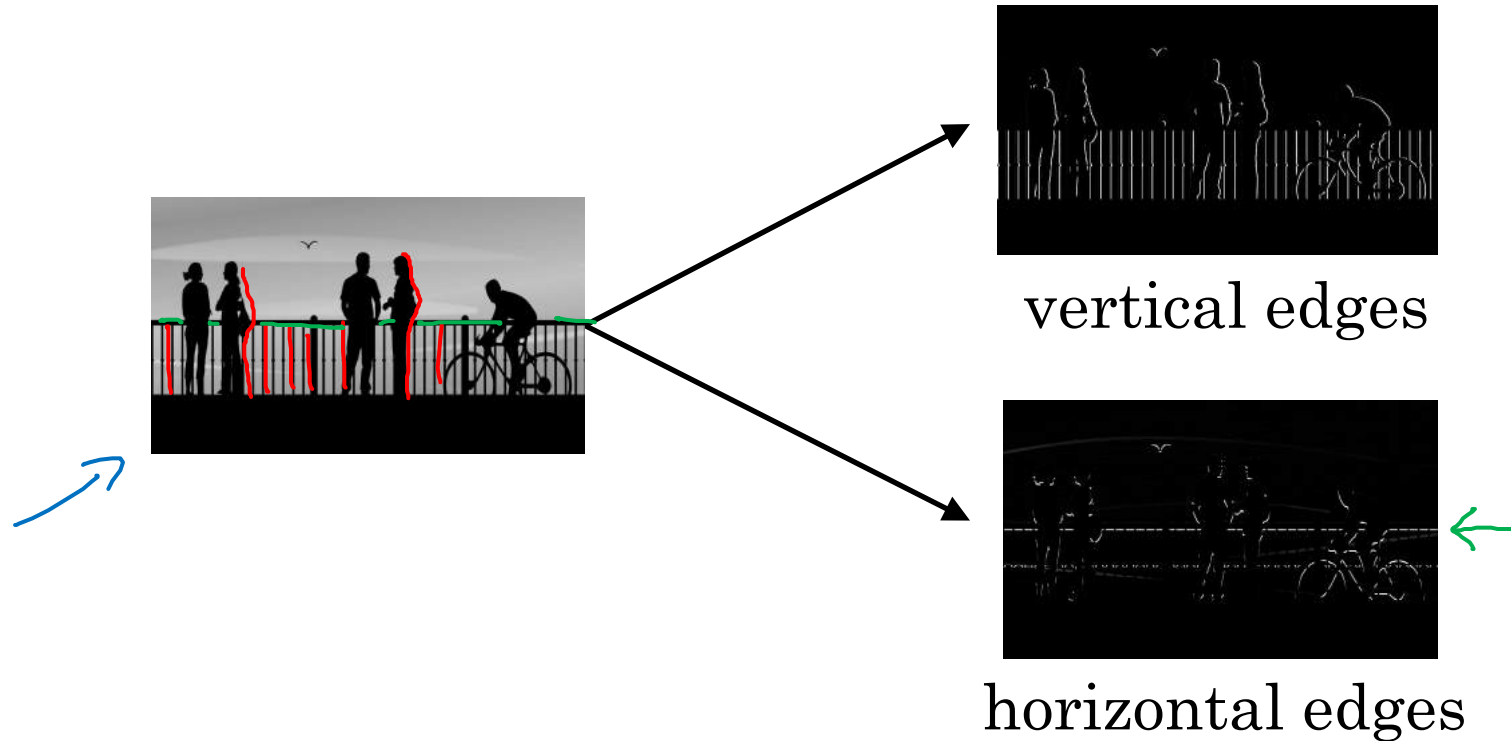
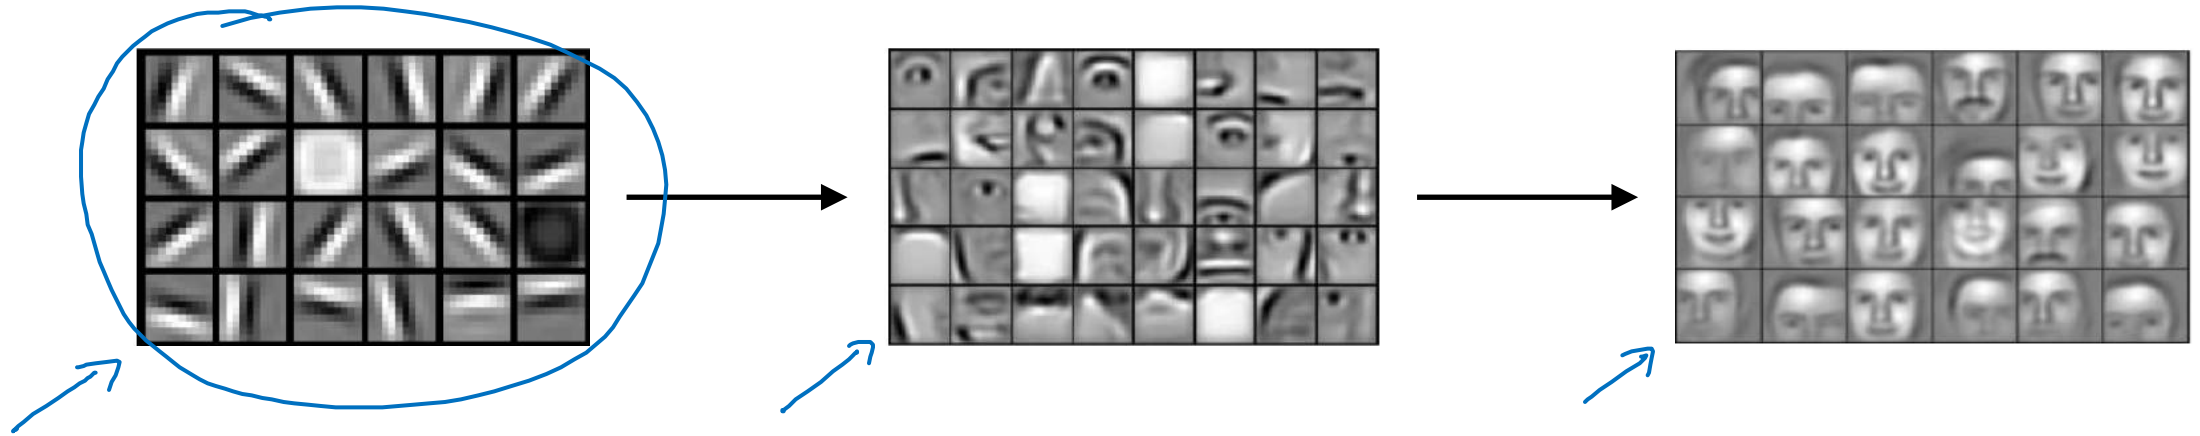


deeplearning.ai

Convolutional Neural Networks

Edge detection example

Computer Vision Problem




Vertical edge detection

$$\rightarrow 3 \times 1 + 1 \times 1 + 2 \times 1 + 0 \times 0 + 5 \times 0 + 7 \times 0 + 1 \times -1 + 8 \times -1 + 2 \times -1 = -5$$

3 ¹	0 ¹	1 ⁻¹	2 ⁻¹	7 ⁻⁰	4 ⁻¹
1 ¹	5 ¹	8 ⁻¹	9 ⁻¹	3 ⁻⁰	1 ⁻¹
2 ¹	7 ¹	2 ⁻¹	5 ⁻¹	1 ⁻⁰	3 ⁻¹
0 ¹	1 ¹	3 ⁻¹	1 ⁻¹	7 ⁻⁰	8 ⁻¹
4	2	1	6	2	8
2	4	5	2	3	9

6x6



=

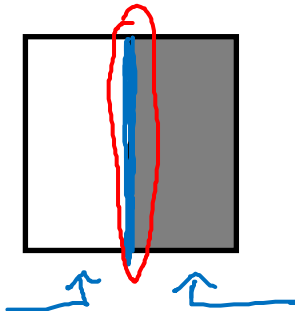
-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

4x4

Vertical edge detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0

6x6

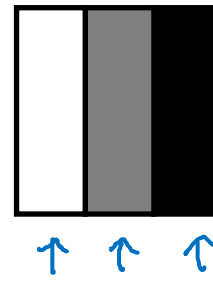


*

1	0	-1
1	0	-1
1	0	-1

3x3

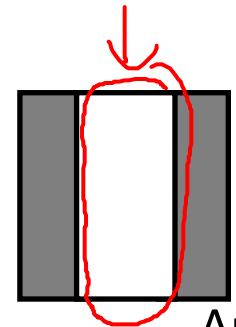
*



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

4x4





deeplearning.ai

Convolutional Neural Networks

More edge
detection

Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



*

1	0	-1
1	0	-1
1	0	-1



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10



*

1	0	-1
1	0	-1
1	0	-1




=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0




Vertical and Horizontal Edge Detection



1	0	-1
1	0	-1
1	0	-1

Vertical



1	1	1
0	0	0
-1	-1	-1

Horizontal

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

6x6

*



1	1	1
0	0	0
-1	-1	-1

=

0	0	0	0
30	10	-10	-30
30	10	-10	-30
0	0	0	0

Learning to detect edges

1	0	-1
1	0	-1
1	0	-1



1	0	-1
2	0	-2
1	0	-1



Sobel filter

3	0	-3
10	0	-10
3	0	-3

Scharr filter



3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

convolution
×

W_1	W_2	W_3
W_4	W_5	W_6
W_7	W_8	W_9

3x3

=

45°
70°
73°



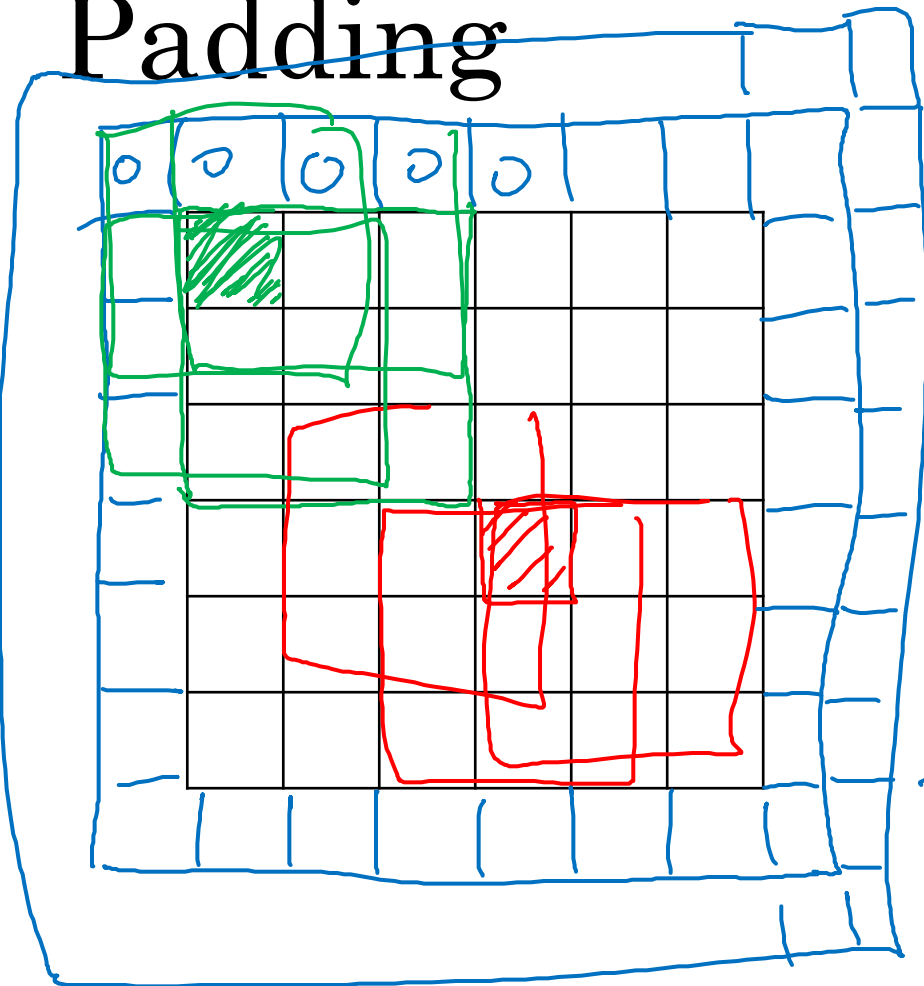
deeplearning.ai

Convolutional Neural Networks

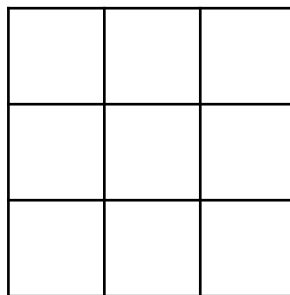
Padding

Padding


- shrinky output
- throw away info from edge



*



=



3×3
 $f \times f$

$p=2$

$\underline{6 \times 6} \rightarrow 8 \times 8$
 $n \times n$

$n - f + 1 \times n - f + 1$
 $6 - 3 + 1 = 4$

$p = \text{padding} = \underline{1}$

$n + 2p - f + 1 \times n + 2p - f + 1$
 $6 + 2 - 3 + 1 \times \underline{\underline{4}} = 6 \times 6$

Valid and Same convolutions

→ no padding

“Valid”: $n \times n \quad * \quad f \times f \quad \rightarrow \quad \frac{n-f+1}{1} \times n-f+1$

$6 \times 6 \quad * \quad 3 \times 3 \quad \rightarrow \quad 4 \times 4$

“Same”: Pad so that output size is the same as the input size.

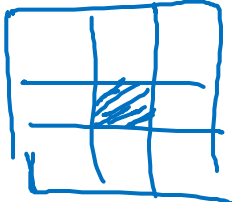
$$n+2p-f+1 \times n+2p-f+1$$
$$\cancel{n+2p-f+1} = \cancel{n} \Rightarrow \boxed{p = \frac{f-1}{2}}$$

$3 \times 3 \quad p = \frac{3-1}{2} = 1$

$5 \times 5 \quad f=5$

f is usually odd

1×1
 3×3
 5×5
 7×7



$p=2$

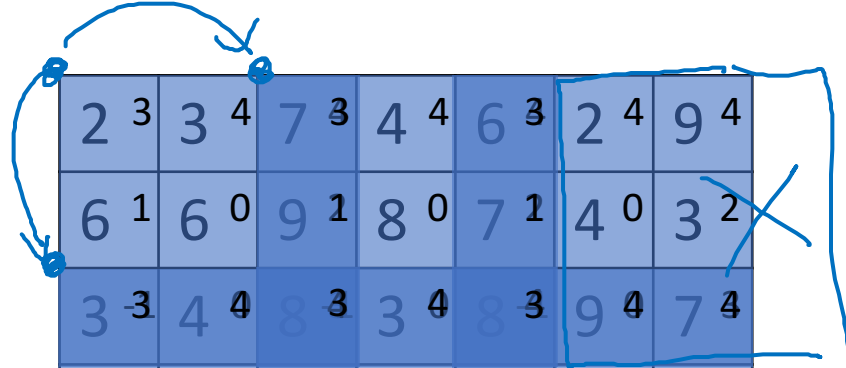


deeplearning.ai

Convolutional Neural Networks

Strided convolutions

Strided convolution



2	3	3	4	7	3	4	4	6	3	2	4	9	4
6	1	6	0	9	1	8	0	7	1	4	0	3	2
3	3	4	4	8	3	3	4	8	3	9	4	7	4
7	1	8	0	3	1	6	0	6	1	3	0	4	2
4	3	2	4	1	3	8	4	3	3	4	4	6	4
3	1	2	0	4	1	1	0	9	1	8	0	3	2
0	-1	1	0	3	-1	9	0	2	-1	1	0	4	3

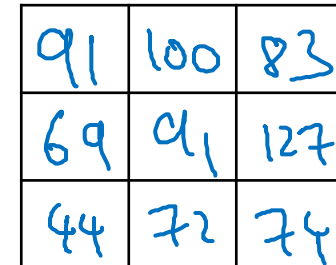
7x7

*

3	4	4
1	0	2
-1	0	3

3x3

=



91	100	83
69	91	127
44	72	74

3x3

Stride = 2

$\lfloor z \rfloor = \text{floor}(z)$

$n \times n$ * $f \times f$
 padding p stride s
 $s = 2$

$$\left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$

$$\frac{7 + 0 - 3}{2} + 1 = \frac{4}{2} + 1 = 3$$

Summary of convolutions

$n \times n$ image $f \times f$ filter

padding p stride s

Output Size:

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \underbrace{\frac{n+2p-f}{s}} + 1 \right\rfloor$$


Technical note on cross-correlation vs. convolution

Convolution in math textbook:

2 ⁻⁷	3 ⁻²	7 ⁻⁵	4	6	2
6 ⁻⁹	6 ⁻⁰	9 ⁻⁴	8	7	4
3 ⁻¹	4 ⁻¹	8 ⁻³	3	8	9
7	8	3	6	6	3
4	2	1	8	3	4
3	2	4	1	9	8

3	4	5
1	0	2
-1	9	7

7	9	5
9	0	4
5	4	3



$$(A * B) * C = A * (B * C)$$

The correct filter after flipping vertically and horizontally would be:




deeplearning.ai

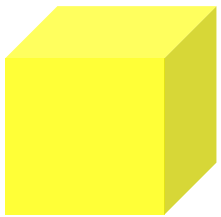
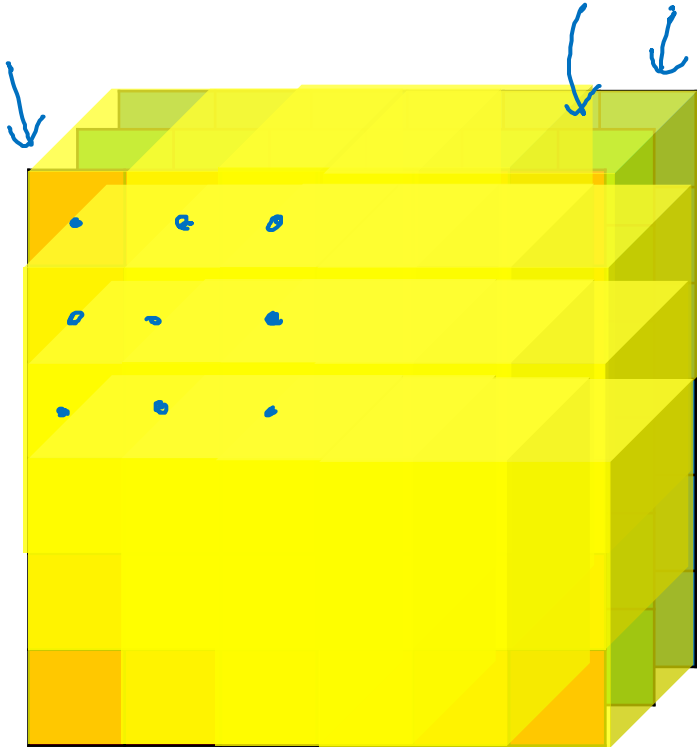
Convolutional Neural Networks

Convolutions over volumes

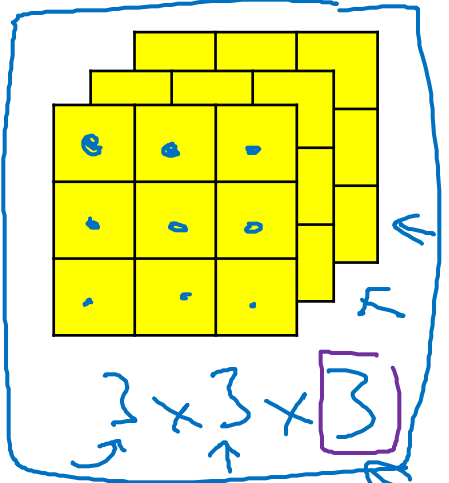
Convolutions on RGB images



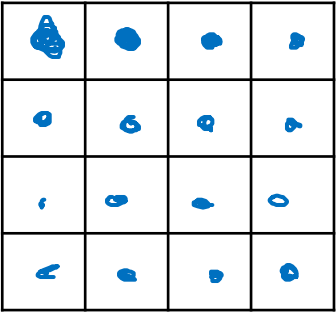
Convolutions on RGB image



*

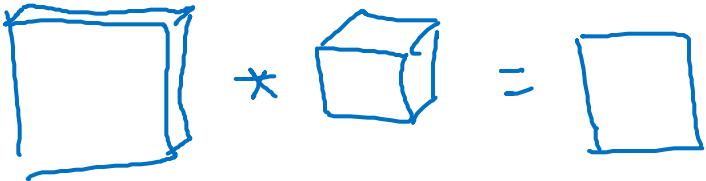


=

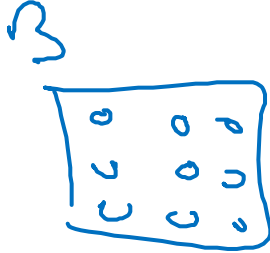
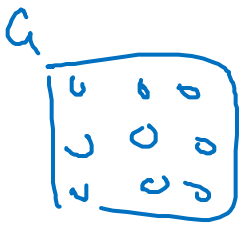
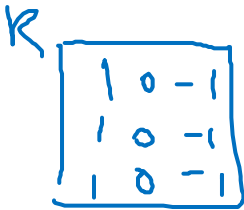


4 x 4

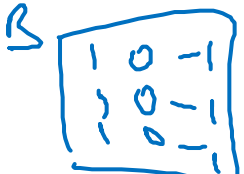
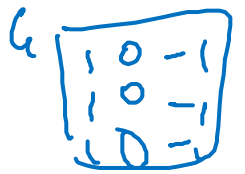
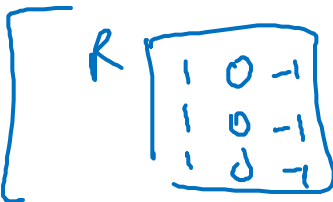
$6 \times 6 \times 3$
↑ ↑ ↑



27 numbers

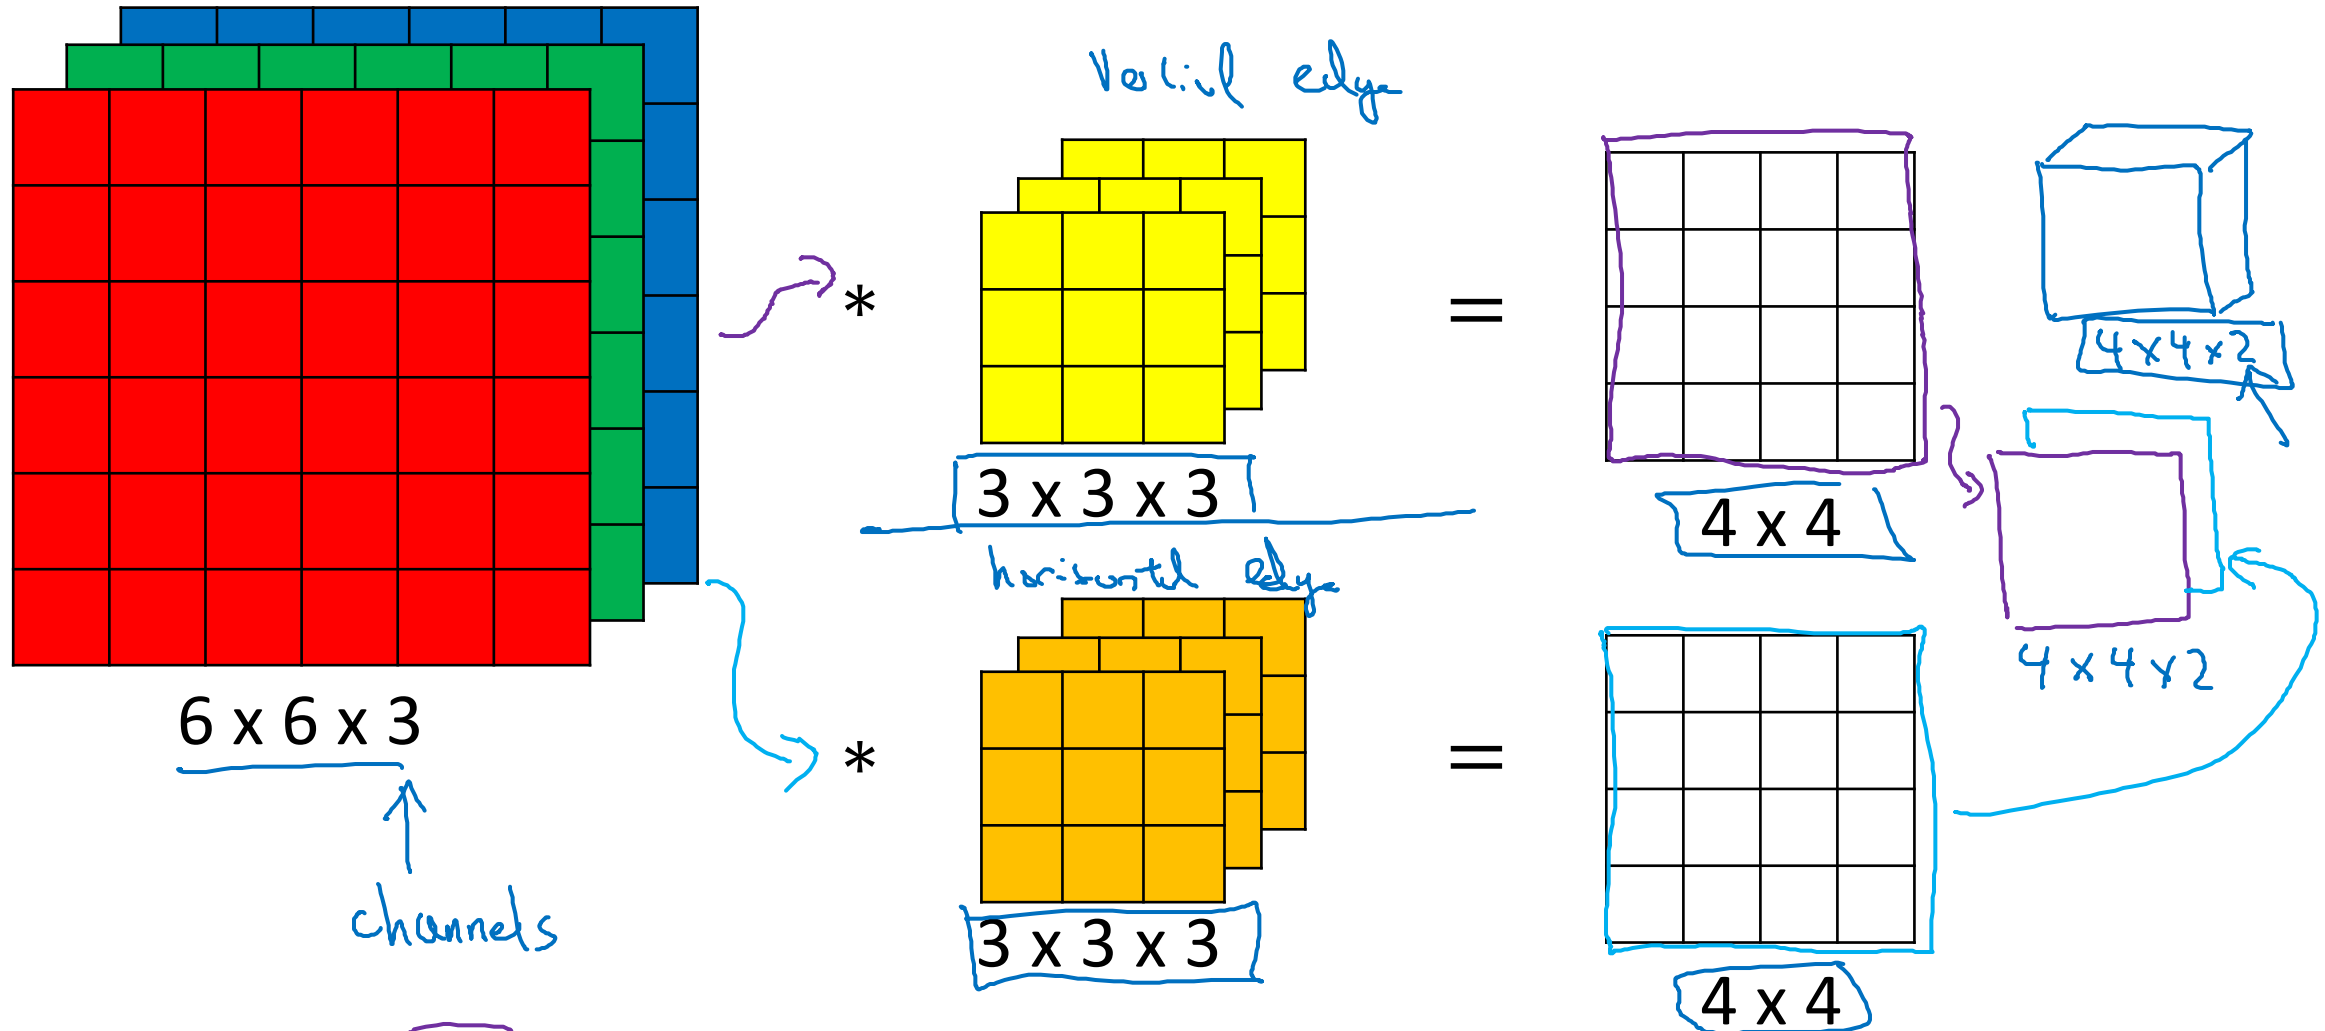


→ 3x3x3



→ 3x3x3

Multiple filters



Summary: $n \times n \times n_c$ \times $f \times f \times n_c$ \rightarrow $\frac{n-f+1}{4} \times \frac{n-f+1}{4} \times n_c'$

$6 \times 6 \times 3$ $3 \times 3 \times 3$ $4 \times 4 \times 2$ \uparrow #filters




deeplearning.ai

Convolutional Neural Networks

One layer of a
convolutional
network

Example of a layer




$6 \times 6 \times 3$

$a^{[0]}$


$$z^{[1]} = W^{[1]} a^{[0]} + b^{[1]}$$

$$a^{[1]} = g(z^{[1]})$$

$*$




$3 \times 3 \times 3$



$3 \times 3 \times 3$

" $W^{[1]}$ "

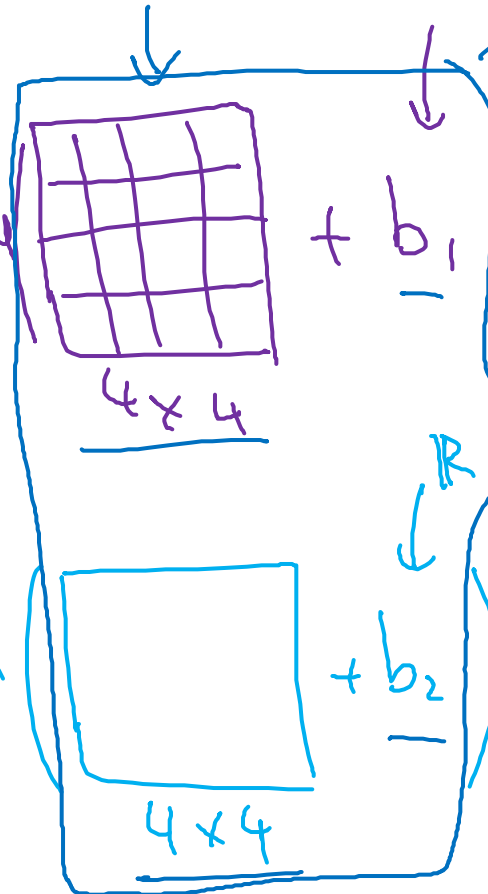


$6 \times 6 \times 3$


$a^{[0]}$

$\rightarrow \text{ReLU}$

$\rightarrow \text{ReLU}$




\rightarrow




4×4

\rightarrow



4×4

\rightarrow



$4 \times 4 \times 2$

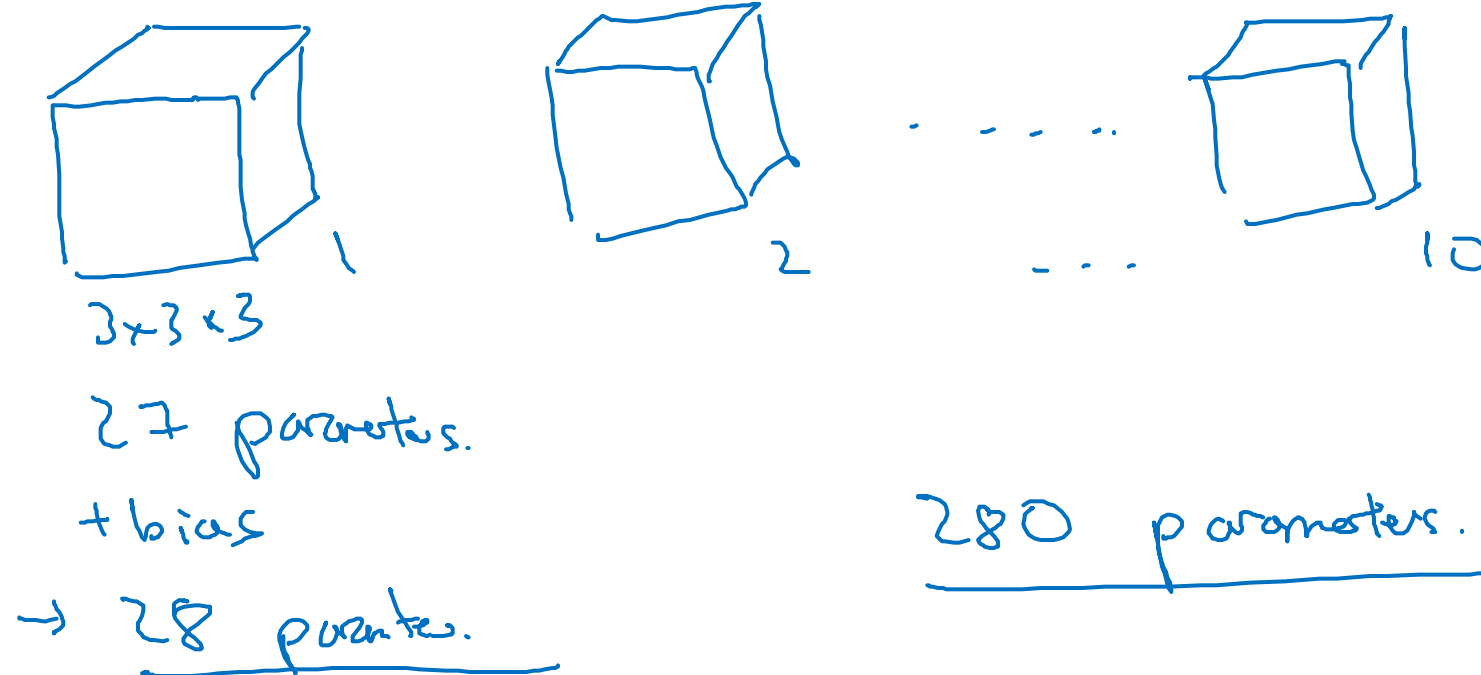
\leftarrow

$a^{[1]}$

$4 \times 4 \times 10$

Number of parameters in one layer

If you have 10 filters that are $3 \times 3 \times 3$ in one layer of a neural network, how many parameters does that layer have?



Summary of notation ★

If layer l is a convolution layer: height, width, channel

$f^{[l]}$ = filter size

$p^{[l]}$ = padding

$s^{[l]}$ = stride

$n_c^{[l]}$ = number of filters

→ Each filter is: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]}$

Activations: $a^{[l]} \rightarrow n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$

Weights: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]}$

bias: $n_c^{[l]} - (1, 1, 1, n_c^{[l]})$ ← #filters in layer l.

Input:

Output:

$$n_{HW}^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rfloor \quad \text{floor}$$

$$A^{[l]} \rightarrow m \times \underbrace{n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}}_{n_c^{[l]} \times n_H^{[l]} \times n_W^{[l]}}$$




deeplearning.ai

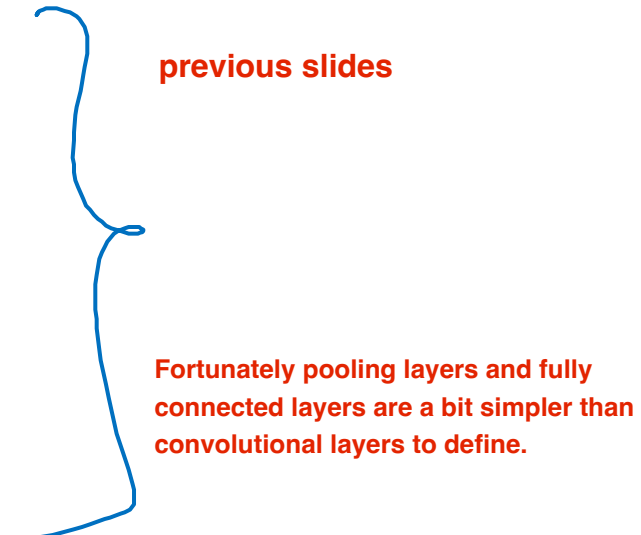
Convolutional Neural Networks

A simple convolution
network example

Example ConvNet



Types of layer in a convolutional network:

- Convolution (conv) ←
 - Pooling (pool) ←
 - Fully connected (FC) ←
- 
- previous slides
- Fortunately pooling layers and fully connected layers are a bit simpler than convolutional layers to define.



deeplearning.ai

Convolutional Neural Networks

Pooling layers

Pooling layer: Max pooling

1	3	2	1
2	9	1	1
1	3	2	3
5	6	1	2

4×4



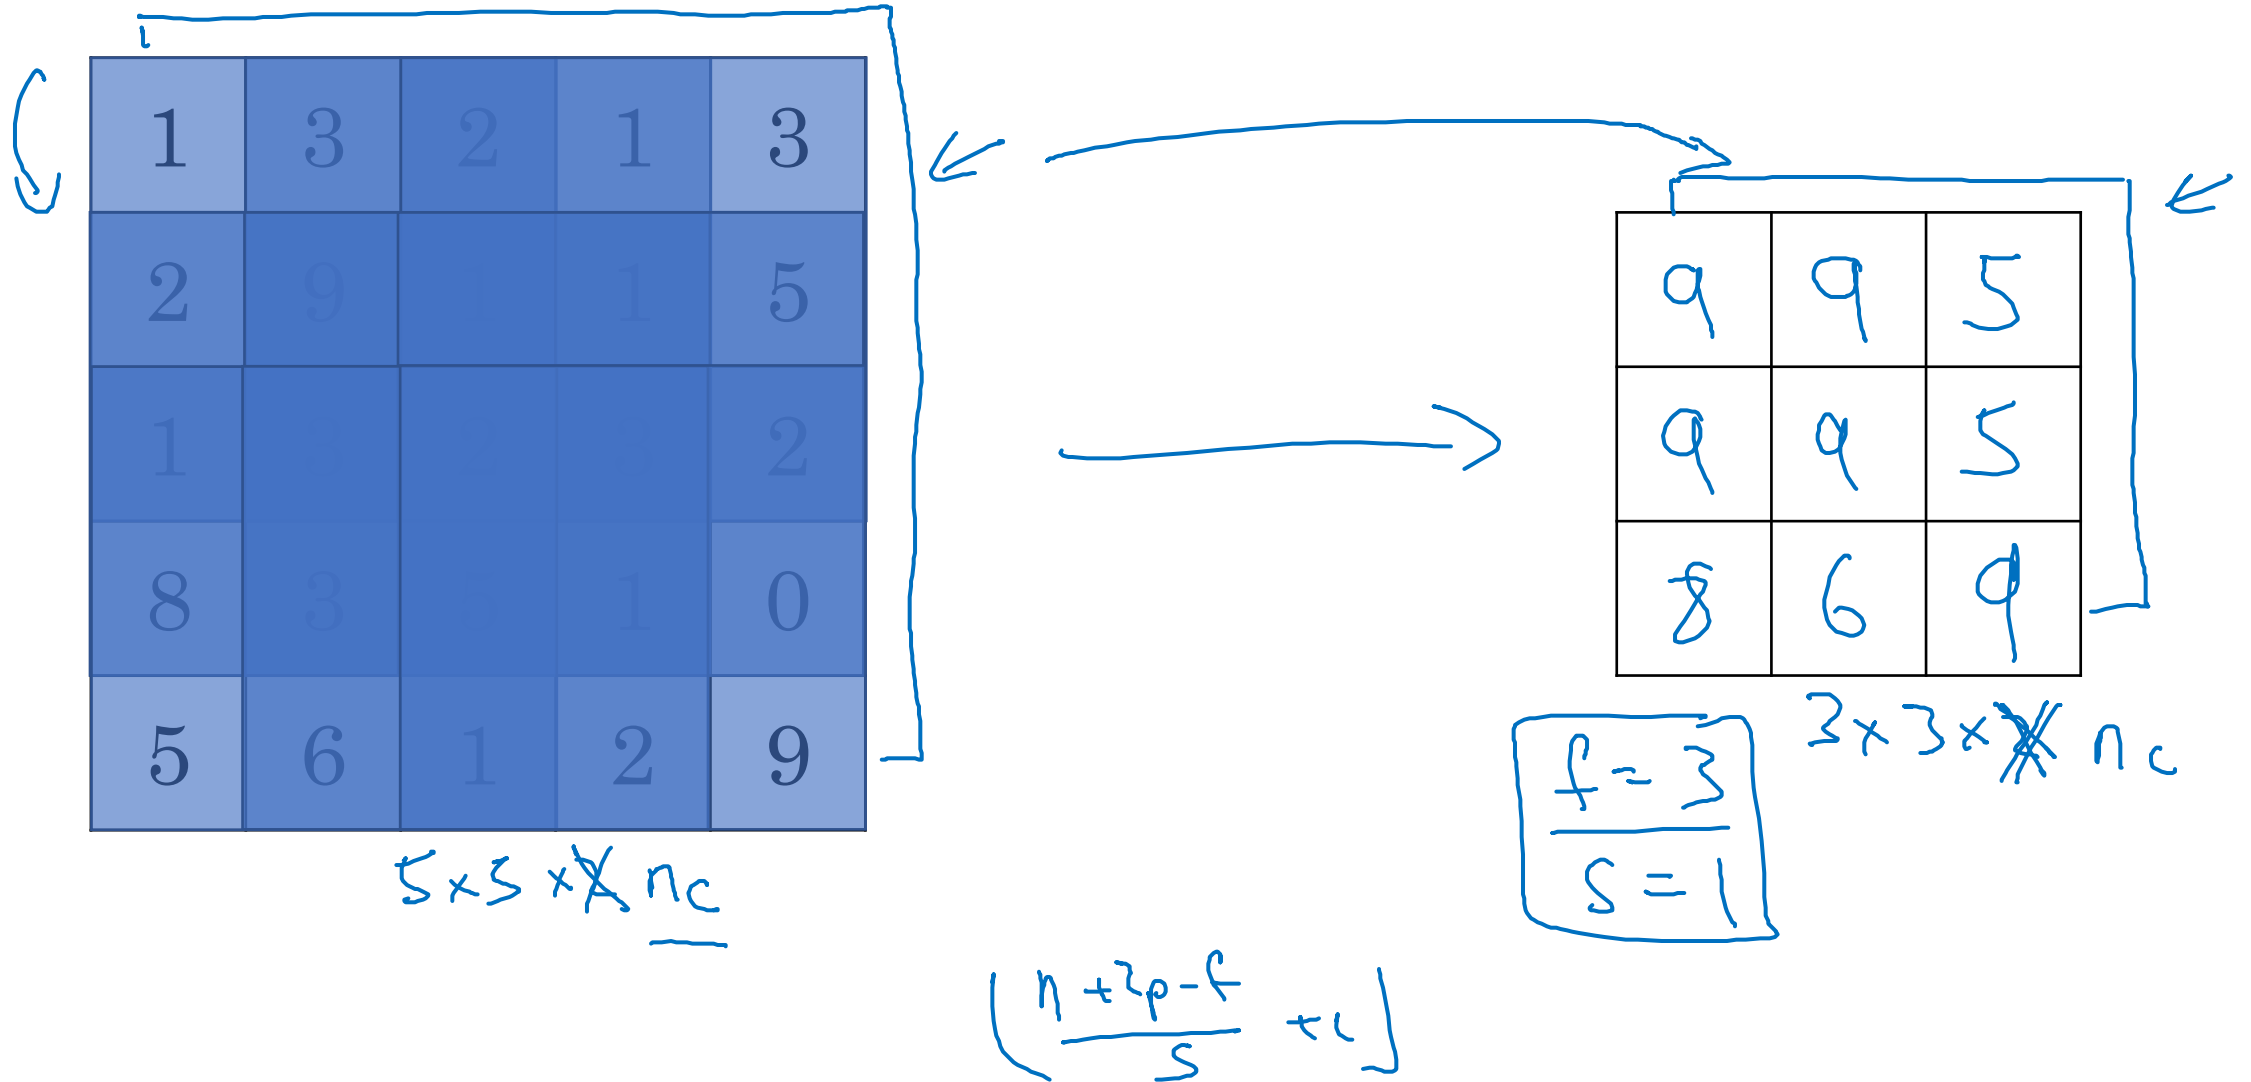
Hyperparameters:
filter=2
stride=2

9	2
6	3

2×2

One interesting property of max pooling is that it has a set of hyperparameters but it has no parameters to learn. There's actually nothing for gradient descent to learn. Once you fix f and s , it's just a fixed computation and gradient descent doesn't change anything.

Pooling layer: Max pooling



Pooling layer: Average pooling

1	3	2	1
2	9	1	1
1	4	2	3
5	6	1	2



3.75	1.25
4	2

$$f=2$$

$$s=2$$

$$\underline{7 \times 7 \times 1000} \rightarrow 1 \times 1 \times 1000$$

max pooling is used much more often than average pooling with one exception, which is sometimes very deep in a neural network.

Summary of pooling

Hyperparameters:

f : filter size

s : stride

Max or average pooling

When you do max pooling, usually, you do not use any padding, although there is one exception that we'll see next week as well.

~~→ p: padding.~~

No parameters to learn!

It's just a fixed function.

$$n_H \times n_W \times \underline{n_C}$$

$$\downarrow$$
$$\left\lfloor \frac{n_H - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n_W - f}{s} + 1 \right\rfloor$$
$$\times \underline{n_C}$$

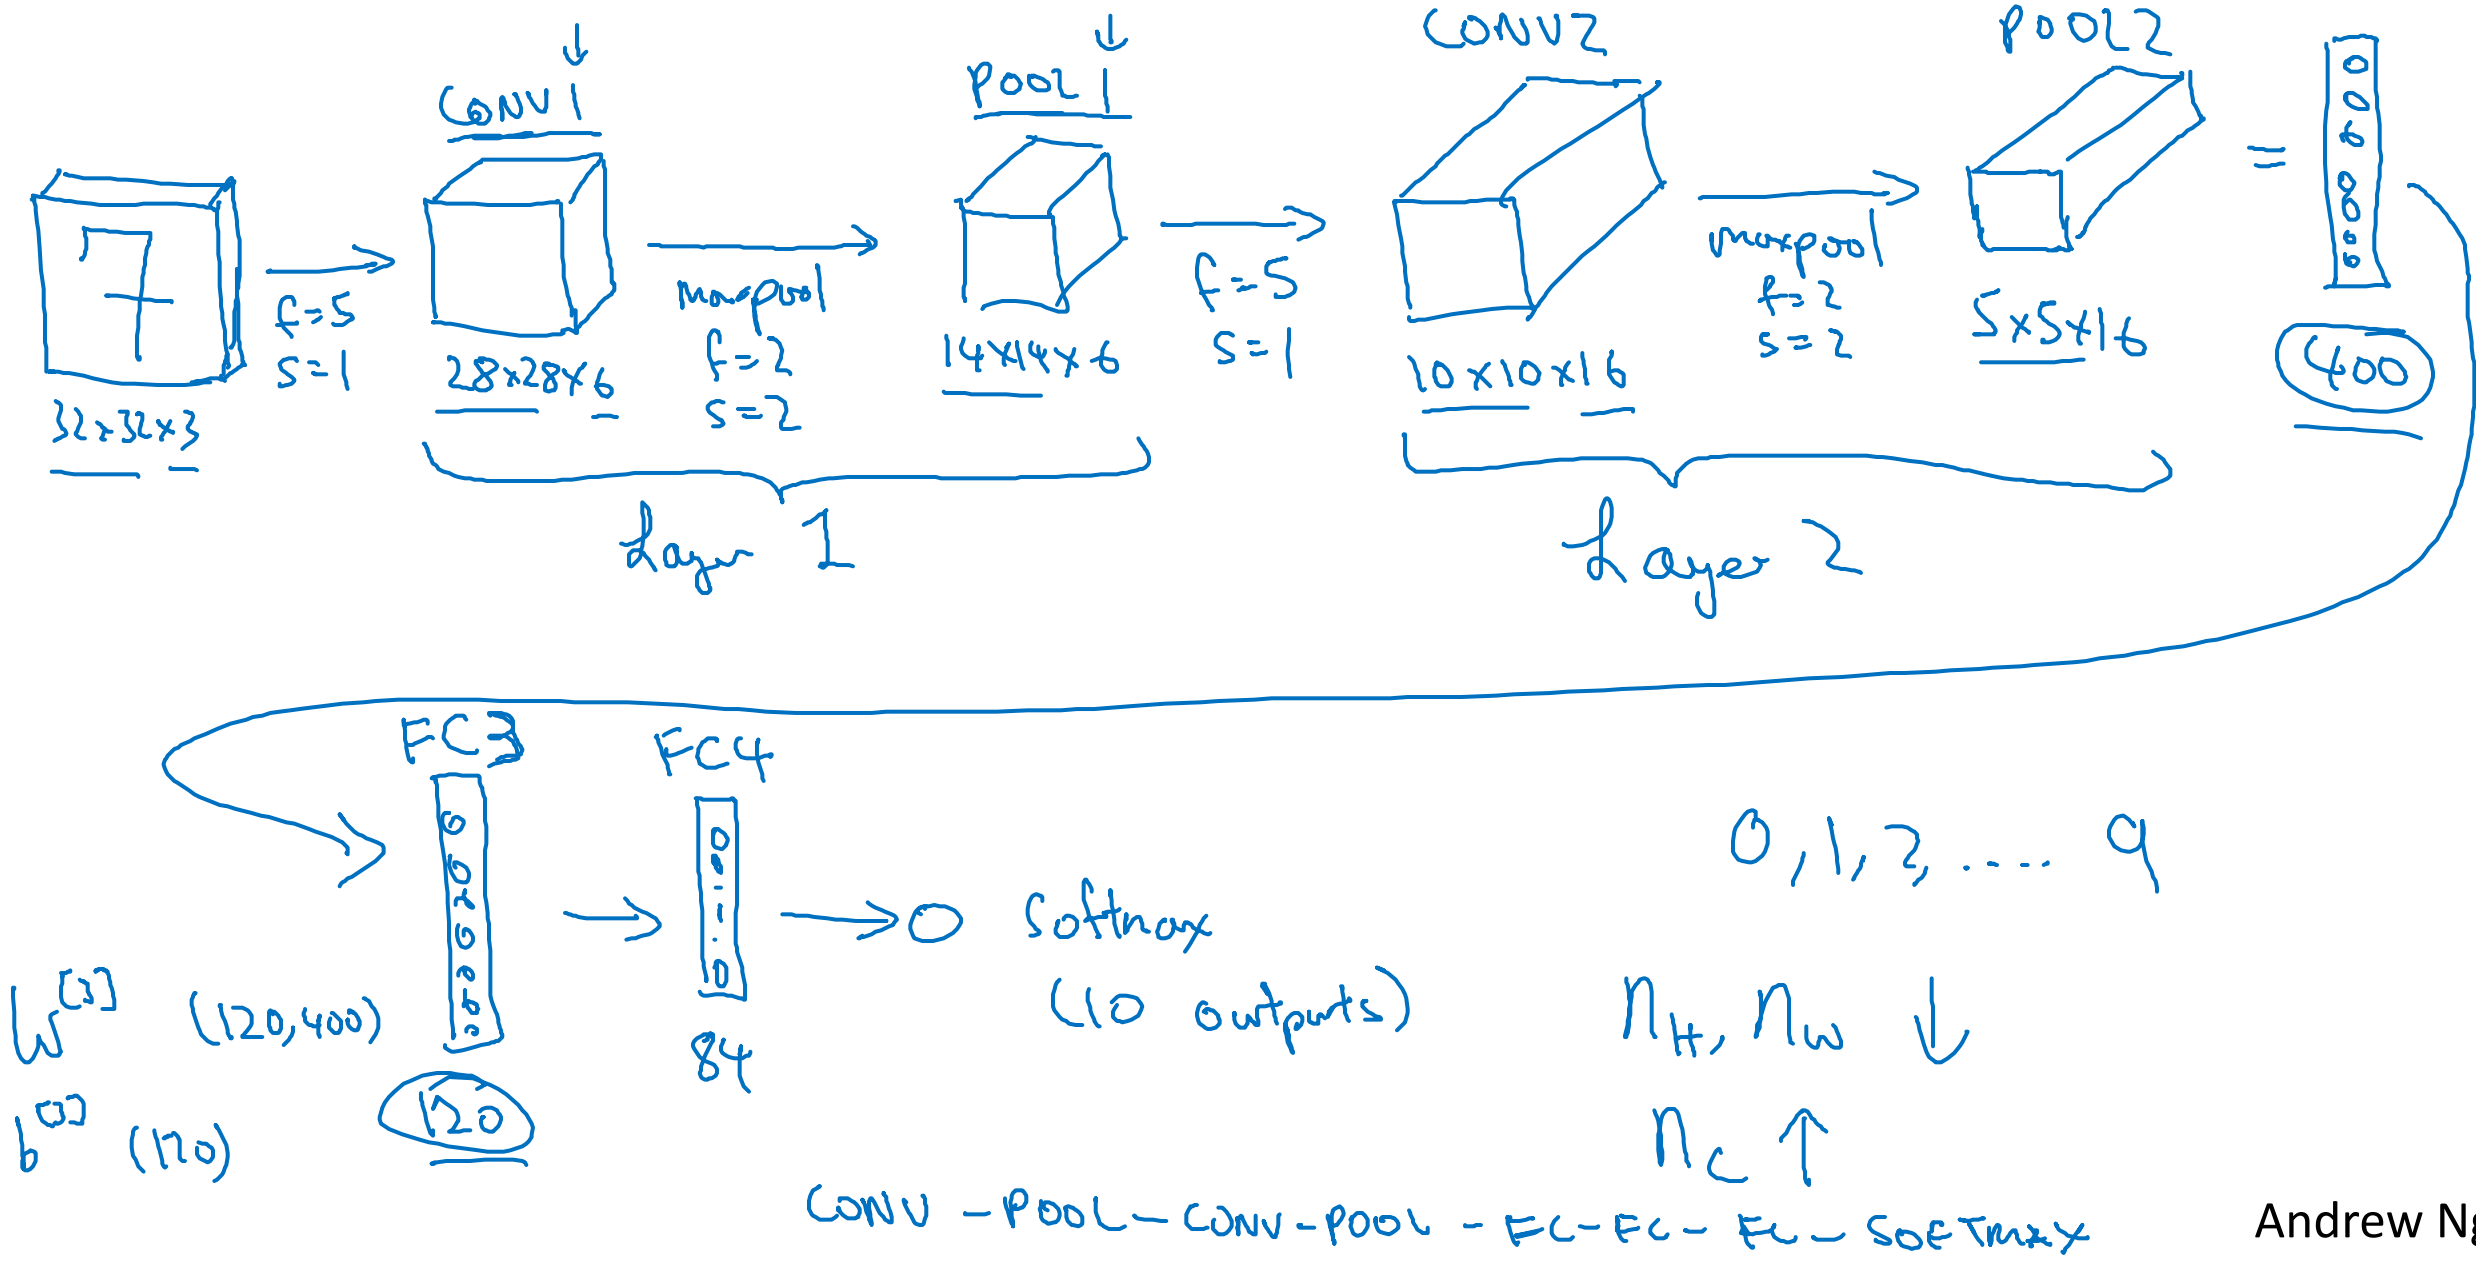


deeplearning.ai

Convolutional Neural Networks

Convolutional neural network example

Neural network example (LeNet-5)



Neural network example

	Activation shape	Activation Size	# parameters
Input:	(32,32,3)	3,072 $a^{[0]}$	0
CONV1 (f=5, s=1)	(28,28,8)	<u>6,272</u>	608 \leftarrow <small>$(5*5*3 + 1) * 8 = 608$</small>
POOL1	(14,14,8)	<u>1,568</u>	0 \leftarrow
CONV2 (f=5, s=1)	(10,10,16)	<u>1,600</u>	3216 \leftarrow
POOL2	(5,5,16)	<u>400</u>	0 \leftarrow
FC3	(120,1)	<u>120</u>	48120 $\left. \vphantom{\begin{matrix} 48120 \\ 10164 \end{matrix}} \right\}$
FC4	(84,1)	<u>84</u>	10164 $\left. \vphantom{\begin{matrix} 48120 \\ 10164 \end{matrix}} \right\}$
Softmax	(10,1)	<u>10</u>	850

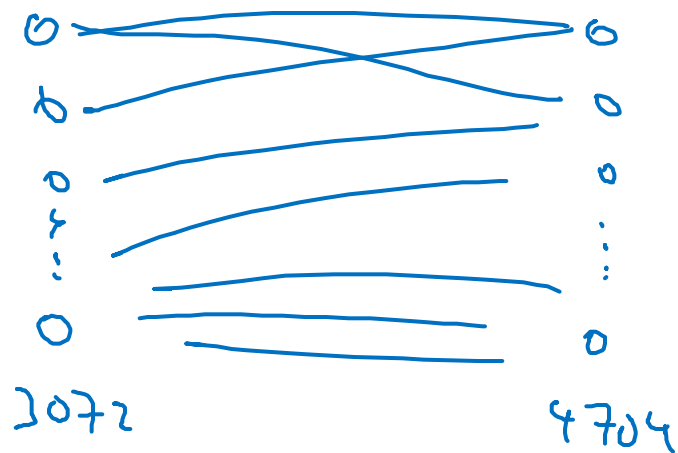
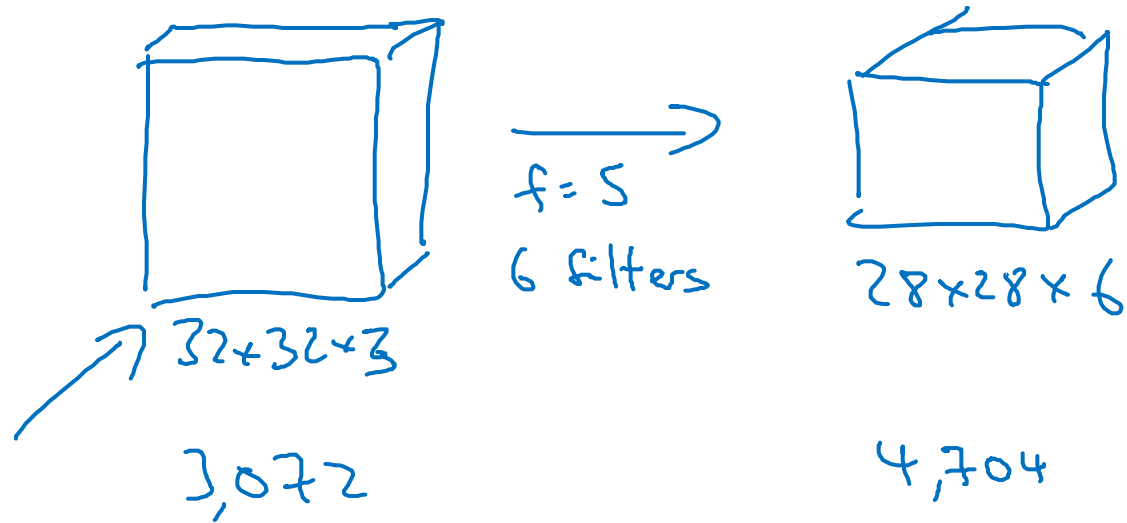


deeplearning.ai

Convolutional Neural Networks

Why convolutions?

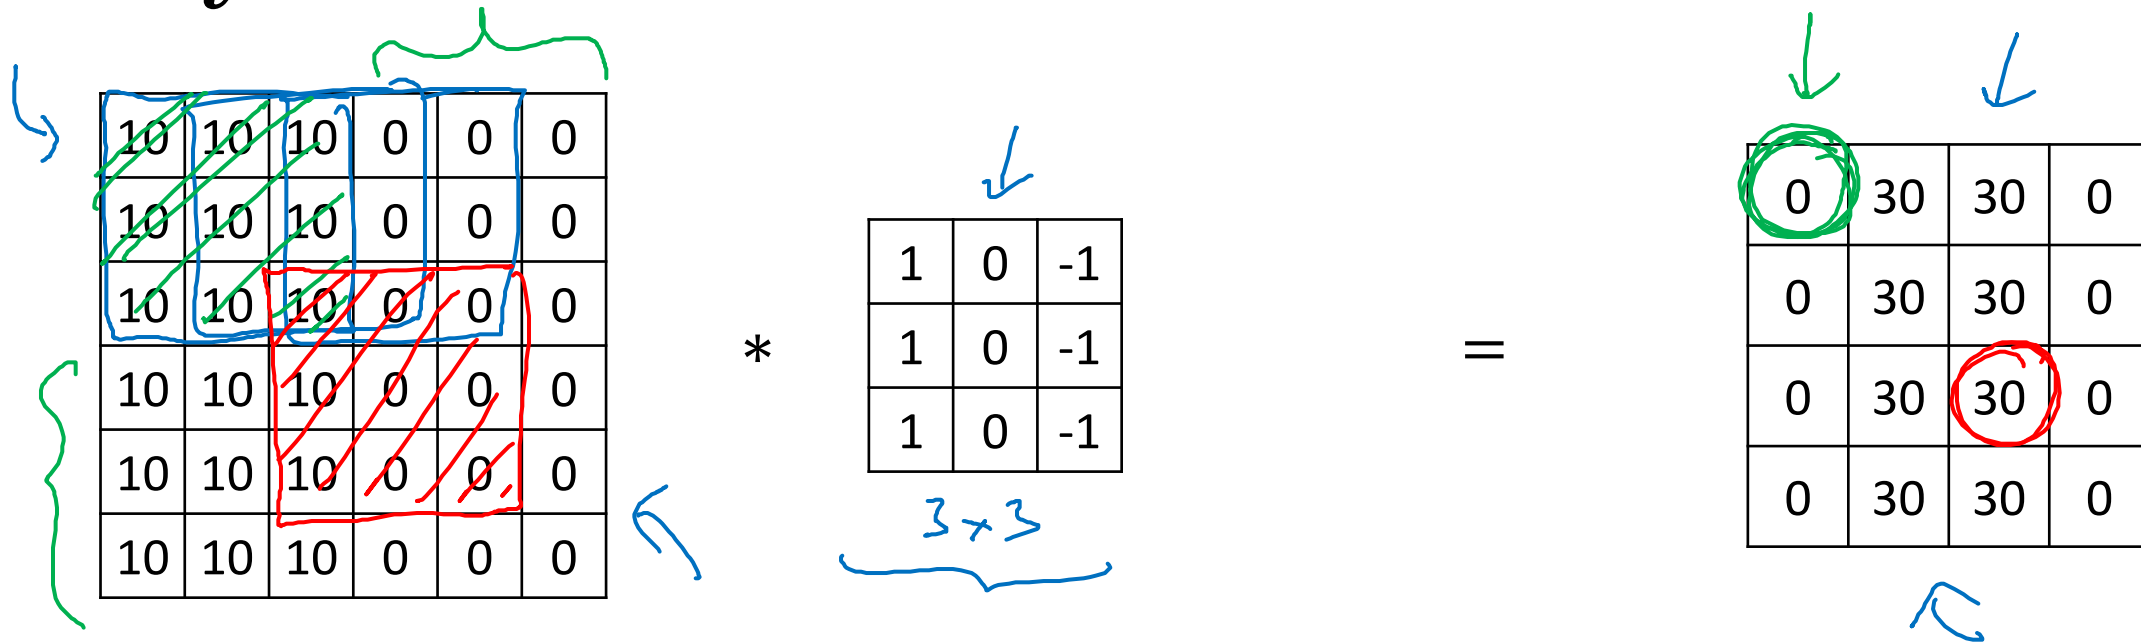
Why convolutions



$$5 \times 5 = 25$$
$$26$$
$$\cancel{6 \times 26 = 156} \text{ parameters}$$
$$(5 \times 5 \times 3 + 1) \times 6 = 456$$

$$3,072 \times 4,704 \approx \underline{14M}$$

Why convolutions

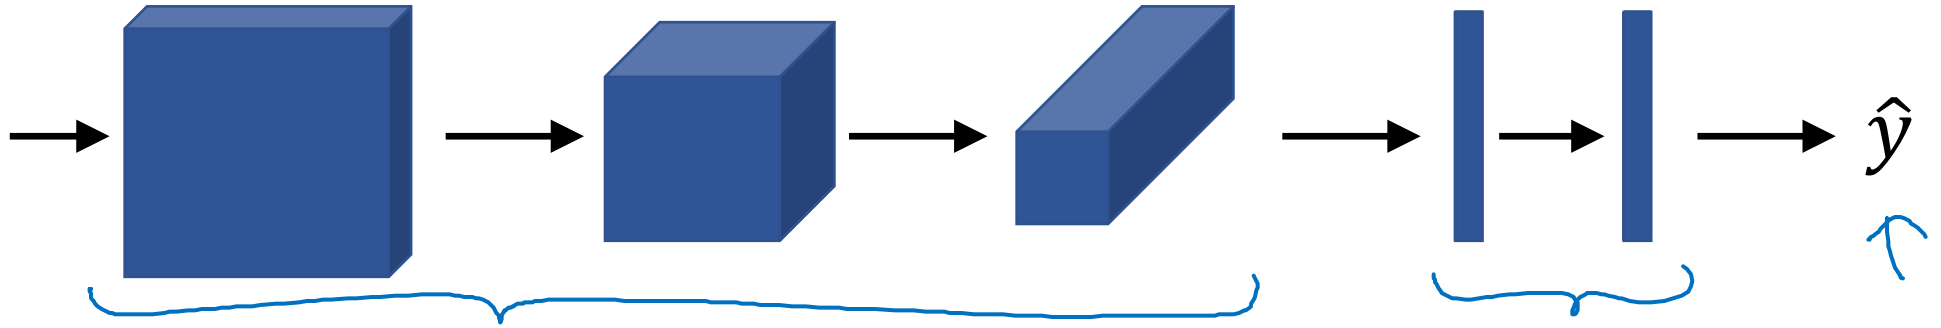


Parameter sharing: A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

→ **Sparsity of connections:** In each layer, each output value depends only on a small number of inputs.

Putting it together

Training set $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$.



$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Use gradient descent to optimize parameters to reduce J